

```
In [1]: import pandas as pd

df = pd.read_csv("water_potability.csv") # or the exact name you see
print(df.head())
print(df.info())
print(df["Potability"].value_counts())
```

	ph	Hardness	Solids	Chloramines	Sulfate	Cond
0	NaN	204.890455	20791.318981	7.300212	368.516441	56
1	3.716080	129.422921	18630.057858	6.635246	NaN	59
2	8.099124	224.236259	19909.541732	9.275884	NaN	41
3	8.316766	214.373394	22018.417441	8.059332	356.886136	36
4	9.092223	181.101509	17978.986339	6.546600	310.135738	39

	Organic_carbon	Trihalomethanes	Turbidity	Potability
0	10.379783	86.990970	2.963135	0
1	15.180013	56.329076	4.500656	0
2	16.868637	66.420093	3.055934	0
3	18.436524	100.341674	4.628771	0
4	11.558279	31.997993	4.075075	0

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3276 entries, 0 to 3275
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ph               2785 non-null    float64
 1   Hardness         3276 non-null    float64
 2   Solids          3276 non-null    float64
 3   Chloramines     3276 non-null    float64
 4   Sulfate          2495 non-null    float64
 5   Conductivity    3276 non-null    float64
 6   Organic_carbon  3276 non-null    float64
 7   Trihalomethanes 3114 non-null    float64
 8   Turbidity        3276 non-null    float64
 9   Potability       3276 non-null    int64  
dtypes: float64(9), int64(1)
memory usage: 256.1 KB
None
Potability
0    1998
1    1278
Name: count, dtype: int64
```

```
In [2]: # 1. Imports
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.metrics import (
    accuracy_score, classification_report, confusion_matrix, roc_auc
)

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
import joblib

# 2. Load your CSV
df = pd.read_csv("water_potability.csv") # put your exact filename here
print("Head:\n", df.head())
print("\nInfo:\n")
print(df.info())
print("\nTarget counts:\n", df["Potability"].value_counts())

# 3. EDA
df.hist(bins=30, figsize=(12, 8))
plt.tight_layout()
plt.show()

plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=False, cmap="coolwarm")
plt.title("Correlation Heatmap")
plt.show()

# 4. Features and target
X = df.drop("Potability", axis=1)
y = df["Potability"]

# 5. Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

# 6. Preprocessing: impute + scale
numeric_features = X.columns.tolist()
numeric_transformer = Pipeline(steps=[
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())
])
preprocessor = ColumnTransformer(
    transformers=[("num", numeric_transformer, numeric_features)]
)

# 7. Models
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Random Forest": RandomForestClassifier(n_estimators=200, random_state=42)
}
```

```
"SVM": SVC(probability=True, random_state=42)
}

results = []

# 8. Train, evaluate, plot
for name, clf in models.items():
    model = Pipeline(steps=[("preprocess", preprocess),
                           ("clf", clf)])
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)
    y_proba = model.predict_proba(X_test)[:, 1]

    acc = accuracy_score(y_test, y_pred)
    roc = roc_auc_score(y_test, y_proba)

    print(f"\n==== {name} ===")
    print("Accuracy:", acc)
    print("ROC-AUC:", roc)
    print("Classification report:")
    print(classification_report(y_test, y_pred))

    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(4, 3))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
    plt.title(f"{name} - Confusion Matrix")
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.show()

    fpr, tpr, _ = roc_curve(y_test, y_proba)
    plt.figure(figsize=(4, 3))
    plt.plot(fpr, tpr, label=f"{name} (AUC={roc:.2f})")
    plt.plot([0, 1], [0, 1], "k--")
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.title(f"{name} - ROC Curve")
    plt.legend()
    plt.show()

    results.append({"Model": name, "Accuracy": acc, "ROC_AUC": roc})

# 9. Compare models
results_df = pd.DataFrame(results).sort_values(by="Accuracy", ascending=False)
print("\nModel performance:\n", results_df)

# 10. Train and save best model (Random Forest example)
best_model = Pipeline(steps=[("preprocess", preprocess),
                           ("clf", RandomForestClassifier(
                               n_estimators=200, random_state=42
                           ))])
best_model.fit(X, y)
joblib.dump(best_model, "water_potability_model.pkl")
print("\nSaved best model as water_potability_model.pkl")
```

Head:

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity
0	NaN	204.890455	20791.318981	7.300212	368.516441	56
1	4.308654	3.716080	129.422921	18630.057858	6.635246	2.885359
2	8.099124	224.236259	19909.541732	9.275884	NaN	8.606213
3	8.316766	214.373394	22018.417441	8.059332	356.886136	3.266516
4	9.092223	181.101509	17978.986339	6.546600	310.135738	8.410813

	Organic_carbon	Trihalomethanes	Turbidity	Potability
0	10.379783	86.990970	2.963135	0
1	15.180013	56.329076	4.500656	0
2	16.868637	66.420093	3.055934	0
3	18.436524	100.341674	4.628771	0
4	11.558279	31.997993	4.075075	0



Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3276 entries, 0 to 3275
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ph               2785 non-null    float64
 1   Hardness          3276 non-null    float64
 2   Solids            3276 non-null    float64
 3   Chloramines       3276 non-null    float64
 4   Sulfate            2495 non-null    float64
 5   Conductivity      3276 non-null    float64
 6   Organic_carbon    3276 non-null    float64
 7   Trihalomethanes   3114 non-null    float64
 8   Turbidity          3276 non-null    float64
 9   Potability         3276 non-null    int64  
dtypes: float64(9), int64(1)
memory usage: 256.1 KB
None
```

Target counts:

	Potability
0	1998
1	1278

Name: count, dtype: int64



== Logistic Regression ==

Accuracy: 0.6097560975609756

ROC-AUC: 0.548115234375

Classification report:

	precision	recall	f1-score	support
0	0.61	1.00	0.76	400
1	0.00	0.00	0.00	256
accuracy			0.61	656
macro avg	0.30	0.50	0.38	656
weighted avg	0.37	0.61	0.46	656

/Library/Frameworks/Python.framework/Versions/3.14/lib/python3.14/site-packages/sklearn/metrics/_classification.py:1731: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])


/Library/Frameworks/Python.framework/Versions/3.14/lib/python3.14/site-packages/sklearn/metrics/_classification.py:1731: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.


_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])

/Library/Frameworks/Python.framework/Versions/3.14/lib/python3.14/site-packages/sklearn/metrics/_classification.py:1731: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])

Logistic Regression - Confusion Matrix






==== Decision Tree ===


Accuracy: 0.5960365853658537

ROC-AUC: 0.5703125

Classification report:


	precision	recall	f1-score	support
0	0.66	0.69	0.67	400
1	0.48	0.45	0.47	256
accuracy			0.60	656
macro avg	0.57	0.57	0.57	656
weighted avg	0.59	0.60	0.59	656






==== Random Forest ====
Accuracy: 0.6615853658536586
ROC-AUC: 0.6532861328125
Classification report:

	precision	recall	f1-score	support
0	0.67	0.89	0.76	400
1	0.63	0.31	0.42	256
accuracy			0.66	656
macro avg	0.65	0.60	0.59	656
weighted avg	0.66	0.66	0.63	656






==== SVM ===


Accuracy: 0.6707317073170732

ROC-AUC: 0.6487109375

Classification report:

	precision	recall	f1-score	support
0	0.66	0.93	0.77	400
1	0.70	0.27	0.39	256
accuracy			0.67	656
macro avg	0.68	0.60	0.58	656
weighted avg	0.68	0.67	0.62	656





Model performance:

	Model	Accuracy	ROC_AUC
3	SVM	0.670732	0.648711
2	Random Forest	0.661585	0.653286
0	Logistic Regression	0.609756	0.548115
1	Decision Tree	0.596037	0.570312

Saved best model as water_potability_model.pkl

In [2]:

```
import pandas as pd

# 1. load your original csv
df = pd.read_csv("water_potability.csv")

# 2. basic cleaning: fill missing numeric values with column median
df_clean = df.copy()
for col in df_clean.columns:
    if df_clean[col].dtype != "object":
        df_clean[col] = df_clean[col].fillna(df_clean[col].median())

# 3. save clean dataset for submission
df_clean.to_csv("clean_water_potability.csv", index=False)
print("Saved clean_water_potability.csv")
```

Saved clean_water_potability.csv

In []: